

Bandung Institute of Technology

Informatics / School of Electrical Engineering and Informatics

IF5251 - AI in Production

Electricity Demand Forecasting



Group 04

Alif Amirudin	23525048
Fawwaz Abrial Saffa	23525055
Ivan Aldy Ganesen	23525062
Sophia D. Waizer	23525901

Academic Year 2025/2026 - Semester 2

Contents

List of Tables	3
List of Figures	4
1 Requirements	5
1.1 System and Model Goals	5
1.2 User Requirements	6
1.2.1 User Types	6
1.2.2 Use Cases	6
1.3 Environment Assumptions	7
1.4 Metrics for Measurement	8
1.4.1 Functional Metrics	8
1.4.2 Non-Functional Metrics	8
1.5 Failure Mode and Effects Analysis	9
2 Architecture and Design	13
2.1 Modelling Tradeoffs	13
2.2 Deployment Architecture	13
2.3 Data Science Pipeline	13
2.4 Telemetry and Monitoring	14
2.5 Anticipate Evolution	14
2.6 Big Data Processing	14
2.7 Human-AI Design	14
3 Quality Assurance	15
3.1 Model Testing	15
3.2 Data Quality	15
3.3 QA Automation	15
3.4 Testing in Production	15
3.5 Infrastructure Quality	16
3.6 Debugging	16
4 Operations	17
4.1 Unit Testing, Integration Testing and Continuous Training Pipeline	17
4.2 Configuration Management	17
4.3 Monitoring	17
4.4 Versioning	18
4.5 MLOps	18
5 Responsible AI	19

5.1	Privacy	19
5.2	Sensitive Audits	19
5.3	Explainability	19
5.4	Transparency	20

List of Tables

1.3	Use Cases Overview	7
1.6	Non-Functional Metrics Overview	9
1.1	User Requirements Overview	10
1.2	User Requirements Mapping to System Users	11
1.4	Environment Assumptions Overview	12
1.5	Functional Metrics Overview	12
3.1	Model Testing Results Placeholder	15
5.1	Sensitive Audit Checklist Placeholder	19

List of Figures

1.1	Use Case Diagram for Electricity Demand Forecasting System	6
1.2	FMEA Diagram Placeholder	9
2.1	Deployment Architecture Diagram Placeholder	13
2.2	Data Science Pipeline Diagram Placeholder	14
2.3	Big Data Processing Architecture Placeholder	14
4.1	Continuous Training Pipeline Diagram Placeholder	17
4.2	Monitoring Dashboard Placeholder	18
4.3	MLOps Workflow Diagram Placeholder	18
5.1	Model Explainability (SHAP / LIME) Placeholder	19

Chapter 1

Requirements

Electricity demand forecasting is a critical component of modern energy grid management. Accurate demand forecasts allow grid operators to balance supply and load proactively, reduce reliance on expensive peaking generation capacity, and facilitate the integration of variable renewable energy sources. This chapter defines the requirements for an AI-based electricity demand forecasting system that operates across multiple geographic locations, covering its goals, the needs of its intended users, the assumptions about the operating environment, and the metrics used to assess both model and system performance.

1.1 System and Model Goals

The *system goal* is to provide **accurate, automated, and near-real-time electricity demand forecasts** for multiple geographic locations, enabling energy operators and planners to make informed grid management decisions. The system ingests historical hourly consumption data alongside meteorological and temporal inputs spanning from 2016 onwards, and exposes forecast results through a web-based interface.

The *model goal* is to learn a **regression function** that maps a rich feature set which consists of over 30 meteorological variables. These features span from weather features such as temperature, humidity, solar radiation, and wind speed, to location identifiers and temporal features. The target variable, y , represents hourly electricity consumption per location cluster. Concretely, the model is expected to perform the following tasks.

- **Forecast** hourly electricity consumption for up to 24 hours ahead for each monitored location.
- **Generalise** across climatically and geographically diverse locations (e.g. London UK, Oakland CA, Groningen NL, Ottawa CA, Washington DC as the locations currently available).
- **Capture** non-linear relationships between weather drivers and demand, including seasonal patterns, diurnal cycles, and weather extremes.
- **Remain performant** on unseen future data through periodic retraining as new consumption records accumulate.

1.2 User Requirements

Based on domain analysis, the following user requirements have been identified in Table 1.1. These requirements are made on the basis of the intended use cases for the forecasting system, which include operational grid management, short-term energy procurement, and demand response planning.

In addition to that, these requirements assume that the users have a certain level of domain expertise in energy management and are familiar with interpreting time-series forecasts in the context of grid operations.

1.2.1 User Types

The system targets two primary user types: **energy grid operators**, who rely on short-term forecasts for real-time grid balancing decisions, and **energy planning analysts**, who use forecasts to support procurement and demand response planning. Both groups require domain expertise in energy management and are assumed to be comfortable interpreting time-series forecasts in an operational context.

The system is **not intended** for end consumers or non-technical stakeholders. Table 1.2 maps each user requirement to the applicable user type, along with representative usage scenarios and implementation priority.

1.2.2 Use Cases

Based on the user requirements defined earlier, the following use cases have been identified for the electricity demand forecasting system, as illustrated in Figure 1.1.




Figure 1.1: Use Case Diagram for Electricity Demand Forecasting System

Each use case corresponds to a specific interaction between the user and the system, aligned with the requirements outlined in Table 1.1, Table 1.3 explains this further.

Table 1.3: Use Cases Overview

ID	Name	Description	Input	Output	Related Req.
UC-01	Request Forecast	User selects a location and time horizon to request a demand forecast via the web interface.	Location identifier, time horizon (up to 48 h)	Inference forecast result	UR-01, UR-02, UR-03, UR-05, UR-07
UC-02	View Visualisation	User views forecast results as a labelled hourly time-series plot.	Inference forecast result	Time-series plot rendered in the interface	UR-04
UC-03	Validate Forecast	User flags a low-confidence forecast for manual review. Extends UC-01.	Forecast result	Correction queued for retraining	UR-06, UR-08
UC-04	View Model Version	User checks which model version produced a given forecast.	Forecast identifier	Model version and training date	UR-09
UC-05	Fetch Weather Data	System automatically retrieves meteorological data for the requested location. Included in UC-01.	Location identifier, request timestamp	Variable meteorological feature set	UR-03

1.3 Environment Assumptions

With the requirements and goals defined, it is important to explicitly state the assumptions about the operating environment in which the system is designed to function. This is because the assumptions will be the bridge in defining the specifications of the system. Other than that, assumptions will help limit the system's scope and clarify the conditions under which the requirements are expected to hold true. The assumptions which are made for this system is listed in Table 1.4. Deviations from these assumptions may require system or model updates. On top of that, any new assumptions that arises during development or deployment might need a reassessment of the requirements and design choices to ensure continued alignment with user needs and operational realities.

1.4 Metrics for Measurement

Measuring system performance requires a combination of model-level (functional) and system-level (non-functional) metrics. Functional metrics assess the accuracy and reliability of the demand forecasts produced by the regression model. Non-functional metrics assess the operational characteristics of the deployed service, including responsiveness, scalability, and availability.

1.4.1 Functional Metrics

The following functional metrics are used to evaluate the regression model's forecasting accuracy. All metrics are computed on a held-out test set comprising data from a time period not seen during training, across all monitored locations. For each metric, specific target thresholds are defined based on domain standards and operational requirements for electricity demand forecasting, which are shown in Table 1.5.

- **MAE** (Mean Absolute Error): the average absolute deviation between predicted and actual consumption; interpretable in the same unit as the target variable and robust to outliers.
- **RMSE** (Root Mean Squared Error): penalises large forecast errors more heavily than MAE, making it sensitive to extreme deviations that are operationally costly.
- **MAPE** (Mean Absolute Percentage Error): expresses error as a percentage of actual consumption, enabling fair comparison across locations with different absolute demand scales.
- **R²** (Coefficient of Determination): measures the proportion of variance in the target variable explained by the model, where a value of 1.0 indicates a perfect fit.

1.4.2 Non-Functional Metrics

Non-functional metrics govern the quality of the deployed service independent of model accuracy. These metrics ensure the system is operationally reliable for production use by energy grid operators operating on a 24/7 basis. For each non-functional metric, specific target thresholds are defined based on the operational requirements of real-time forecasting services in the energy sector, which are shown in Table 1.6.

- **Latency**: the end-to-end response time from API request submission to forecast delivery. Low latency is critical in grid management contexts where decisions are time-sensitive.
- **Throughput**: the number of concurrent forecast requests the service can handle per unit time without degradation, ensuring the system scales to multiple simultaneous operator queries.
- **Availability**: the percentage of time the forecasting service is accessible. Unplanned downtime has direct operational consequences for grid operators who depend on continuous access.
- **Model Freshness**: the maximum permissible age of the deployed model before retraining is mandated, ensuring that concept drift and seasonal pattern shifts do not degrade forecast accuracy over time.

Table 1.6: Non-Functional Metrics Overview

Metric	Target
Latency	< 500 ms
Throughput	≥ 50 req/s
Availability	$\geq 99.5\%$
Model Freshness	≤ 60 days

1.5 Failure Mode and Effects Analysis

Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Donec odio elit, dictum in, hendrerit sit amet, egestas sed, leo. Praesent feugiat sapien aliquet odio. Integer vitae justo.




Figure 1.2: FMEA Diagram Placeholder

Morbi luctus, wisi viverra faucibus pretium, nibh est placerat odio, nec commodo wisi enim eget quam. Quisque libero justo, consectetur a, feugiat vitae, porttitor eu, libero.

Table 1.1: User Requirements Overview

ID	Requirement Name	Description	Rationale
UR-01	Forecast Query	Users shall be able to request an electricity demand forecast for a specific monitored location and a time horizon of up to 48 hours via a web-based interface.	Enables users to obtain actionable forecasts for operational decision-making.
UR-02	Location Selection	The system shall support all locations present in the training dataset, selectable by name or geographic identifier (e.g. city name or <code>location_id</code>).	Allows users to access forecasts relevant to their specific geographic area of interest.
UR-03	Weather Data Fetch	The system shall fetch the current relevant weather data for each location automatically.	Reduces user burden by eliminating manual weather input, while ensuring forecasts are based on current meteorological data.
UR-04	Forecast Visualisation	Forecast results shall be displayed as a time-series plot of predicted hourly consumption over the selected time window, with clear axis labelling.	Facilitates intuitive understanding of forecast trends and patterns.
UR-05	Response Timeliness	Forecast results shall be returned within an operationally acceptable response time to support real-time grid management decisions.	Ensures forecasts are delivered in a timeframe that allows for timely action.
UR-06	Accuracy Transparency	The interface shall communicate expected model error ranges so users can appropriately calibrate their trust in the forecasts.	Builds user trust by setting realistic expectations about forecast reliability.
UR-07	Continuous Availability	The forecasting service shall be accessible at all times, reflecting the 24/7 nature of grid operations.	Ensures users can access forecasts whenever needed without interruption.
UR-08	Validation Interface	Users must be able to flag or manually validate forecasts that fall below a confidence threshold.	Allows users to ensure forecast quality and accuracy and supplying additional data for retraining.
UR-09	Model Version Visibility	Users must be able to see which version of the forecasting model was used to generate a particular forecast.	Enables users to understand the context and potential staleness of each forecast.

Table 1.2: User Requirements Mapping to System Users

ID	Energy Grid Operator	Energy Planning Analyst	Usage Scenario	Priority
UR-01	✓	✓	Requesting forecasts for operational decision-making.	High
UR-02	✓	✓	Selecting specific locations for forecast retrieval.	High
UR-03	✓	✓	Ensuring queries doesn't burden the user.	High
UR-04	✓	✓	Visualising forecast results for interpretation.	Medium
UR-05	✓	✗	Ensuring forecasts are delivered in a timely manner for real-time actions.	High
UR-06	✓	✓	Understanding expected error ranges to calibrate trust.	Medium
UR-07	✓	✓	Accessing forecasts at any time without interruption.	High
UR-08	✓	✗	Validating forecasts that fall below confidence thresholds.	Medium
UR-09	✓	✓	Seeing which version of the model was used.	Medium.

Table 1.4: Environment Assumptions Overview

ID	Assumption Name	Description
EA-01	Weather Data Availability	Real-time and forecast meteorological data matching the 30+ feature schema used during training is continuously available as model input at inference time.
EA-02	Hourly Temporal Resolution	All historical and incoming data is timestamped at hourly granularity. Sub-hourly or irregular inputs are outside the current scope.
EA-03	Stationarity of Demand Patterns	The relationship between weather conditions, time features, and electricity consumption is assumed to be sufficiently stable over time for the trained model to remain valid between retraining cycles. Major structural disruptions may require earlier retraining.
EA-04	Multi-Location Heterogeneity	Each location has a distinct demand profile driven by local climate, building stock, and population density. The <code>location_id</code> and <code>cluster_size</code> fields capture this heterogeneity and must be present at inference time.
EA-05	Data Quality	Historical data is assumed to have been preprocessed (null values handled, outliers addressed) as reflected in the <code>full_data_non_null.csv</code> dataset. Incoming data at inference time is expected to pass basic schema and range validation checks before being fed to the model.
EA-06	Deployment Infrastructure	The system is deployed on a server or cloud instance capable of hosting the trained model and serving REST API requests from the Next.js frontend.

Table 1.5: Functional Metrics Overview

Metric	Description	Target
MAE	Mean Absolute Error	< 10% of mean target per location
RMSE	Root Mean Squared Error	< 15% of mean target per location
MAPE	Mean Absolute Percentage Error	< 10%
R ²	Coefficient of Determination	> 0.85

Chapter 2

Architecture and Design

Suspendisse vitae elit. Aliquam arcu neque, ornare in, ullamcorper quis, commodo eu, libero. Fusce sagittis erat at erat tristique mollis.

2.1 Modelling Tradeoffs

Sed feugiat. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Ut pellentesque augue sed urna. Vestibulum diam eros, fringilla et, consectetur eu, nonummy id, sapien.

2.2 Deployment Architecture

Etiam euismod. Fusce facilisis lacinia dui. Suspendisse potenti. In mi erat, cursus id, nonummy sed, ullamcorper eget, sapien.




Figure 2.1: Deployment Architecture Diagram Placeholder

Aliquam lectus. Vivamus leo.

2.3 Data Science Pipeline

Etiam ac leo a risus tristique nonummy. Donec dignissim tincidunt nulla. Vestibulum rhoncus molestie odio. Sed lobortis, justo et pretium lobortis, mauris turpis condimentum augue, nec ultricies nibh arcu pretium enim.



Figure 2.2: Data Science Pipeline Diagram Placeholder

Nulla in ipsum. Praesent eros nulla, congue vitae, euismod ut, commodo a, wisi.

2.4 Telemetry and Monitoring

Nulla mattis luctus nulla. Duis commodo velit at leo. Aliquam vulputate magna et leo. Nam vestibulum ullamcorper leo.

2.5 Anticipate Evolution

Curabitur tellus magna, porttitor a, commodo a, commodo in, tortor. Donec interdum. Praesent scelerisque. Maecenas posuere sodales odio.

2.6 Big Data Processing

Donec et nisl at wisi luctus bibendum. Nam interdum tellus ac libero. Sed sem justo, laoreet vitae, fringilla at, adipiscing ut, nibh. Maecenas non sem quis tortor eleifend fermentum.



Figure 2.3: Big Data Processing Architecture Placeholder

Nulla non mauris vitae wisi posuere convallis. Sed eu nulla nec eros scelerisque pharetra.

2.7 Human-AI Design

Nulla ac nisl. Nullam urna nulla, ullamcorper in, interdum sit amet, gravida ut, risus. Aenean ac enim. In luctus.

Chapter 3

Quality Assurance

Etiam pede massa, dapibus vitae, rhoncus in, placerat posuere, odio. Vestibulum luctus commodo lacus. Morbi lacus dui, tempor sed, euismod eget, condimentum at, tortor.

3.1 Model Testing

Etiam suscipit aliquam arcu. Aliquam sit amet est ac purus bibendum congue. Sed in eros. Morbi non orci.

Table 3.1: Model Testing Results Placeholder

Test Case	Expected Output	Actual Output
Case 1	Placeholder	Placeholder
Case 2	Placeholder	Placeholder
Case 3	Placeholder	Placeholder

3.2 Data Quality

Donec et nisl id sapien blandit mattis. Aenean dictum odio sit amet risus. Morbi purus. Nulla a est sit amet purus venenatis iaculis.

3.3 QA Automation

Maecenas non massa. Vestibulum pharetra nulla at lorem. Duis quis quam id lacus dapibus interdum. Nulla lorem.

3.4 Testing in Production

Vivamus eu tellus sed tellus consequat suscipit. Nam orci orci, malesuada id, gravida nec, ultricies vitae, erat. Donec risus turpis, luctus sit amet, interdum quis, porta sed, ipsum.

Adversarial Attacks

Duis aliquet dui in est. Donec eget est. Nunc lectus odio, varius at, fermentum in, accumsan non, enim.

Load Testing

Donec vel nibh ut felis consectetuer laoreet. Donec pede. Sed id quam id wisi laoreet suscipit.

3.5 Infrastructure Quality

Donec molestie, magna ut luctus ultrices, tellus arcu nonummy velit, sit amet pulvinar elit justo et mauris. In pede. Maecenas euismod elit eu erat. Aliquam augue wisi, facilisis congue, suscipit in, adipiscing et, ante.

3.6 Debugging

Cras dapibus, augue quis scelerisque ultricies, felis dolor placerat sem, id porta velit odio eu elit. Aenean interdum nibh sed wisi. Praesent sollicitudin vulputate dui. Praesent iaculis viverra augue.

Chapter 4

Operations

Sed mattis, erat sit amet gravida malesuada, elit augue egestas diam, tempus scelerisque nunc nisl vitae libero. Sed consequat feugiat massa. Nunc porta, eros in eleifend varius, erat leo rutrum dui, non convallis lectus orci ut nibh.

4.1 Unit Testing, Integration Testing and Continuous Training Pipeline

Sed consequat tellus et tortor. Ut tempor laoreet quam. Nullam id wisi a libero tristique semper. Nullam nisl massa, rutrum ut, egestas semper, mollis id, leo.



Figure 4.1: Continuous Training Pipeline Diagram Placeholder

Phasellus id magna. Duis malesuada interdum arcu.

4.2 Configuration Management

Sed eleifend, eros sit amet faucibus elementum, urna sapien consectetur mauris, quis egestas leo justo non risus. Morbi non felis ac libero vulputate fringilla. Mauris libero eros, lacinia non, sodales quis, dapibus porttitor, pede. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos.

4.3 Monitoring

Nullam eleifend justo in nisl. In hac habitasse platea dictumst. Morbi nonummy. Aliquam ut felis.



Figure 4.2: Monitoring Dashboard Placeholder

Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Aenean nonummy turpis id odio.

4.4 Versioning

Nulla malesuada risus ut urna. Aenean pretium velit sit amet metus. Duis iaculis. In hac habitasse platea dictumst.

4.5 MLOps

Donec tempus neque vitae est. Aenean egestas odio sed risus ullamcorper ullamcorper. Sed in nulla a tortor tincidunt egestas. Nam sapien tortor, elementum sit amet, aliquam in, porttitor faucibus, enim.




Figure 4.3: MLOps Workflow Diagram Placeholder

Fusce suscipit cursus sem. Vivamus risus mi, egestas ac, imperdiet varius, faucibus quis, leo.

Chapter 5

Responsible AI

Praesent sed neque id pede mollis rutrum. Vestibulum iaculis risus. Pellentesque lacus.

5.1 Privacy

Sed gravida lectus ut purus. Morbi laoreet magna. Pellentesque eu wisi. Proin turpis.

5.2 Sensitive Audits

Curabitur ac lorem. Vivamus non justo in dui mattis posuere. Etiam accumsan ligula id pede. Maecenas tincidunt diam nec velit.

Table 5.1: Sensitive Audit Checklist Placeholder

Audit Item	Status	Notes
Data Anonymisation	Pending	Placeholder
Bias Assessment	Pending	Placeholder
Fairness Evaluation	Pending	Placeholder

5.3 Explainability

Quisque consectetur. In suscipit mauris a dolor pellentesque consectetur. Mauris convallis neque non erat. In lacinia.



Figure 5.1: Model Explainability (SHAP / LIME) Placeholder

Maecenas accumsan dapibus sapien. Duis pretium iaculis arcu.

5.4 Transparency

Phasellus fringilla, metus id feugiat consectetur, lacus wisi ultrices tellus, quis lobortis nibh lorem quis tortor. Donec egestas ornare nulla. Mauris mi tellus, porta faucibus, dictum vel, nonummy in, est. Aliquam erat volutpat.